

CS143: SQL Query (2)

Book Chapters

- (5th) Chapter 3.5, 3.7
- (6th) Chapter 3.7-8
- (7th) Chapter 3.7-8

Things to Learn

- Subquery
- Aggregate

Subqueries

- **SELECT** statement may appear in **WHERE** clause
 - Treated the same as regular relations
 - If the result is one-attribute one-tuple relation, the result can be used like a 'value'

Scalar-value subqueries

- **Query 1:** Find the student ids who live at the same addr as the student with id 301

- **Q:** Can we rewrite it without subquery?

- **Notes:**

- There is a whole theory about whether/how to rewrite a subquery to non-subquery SQL

- The basic result is we can rewrite subqueries as long as we do not have negation.
- With negation, we need **EXCEPT**
- One of the reasons why relational model has been so successful
 - * Because it is easy to understand and model, we can design and prove elegant theorems.
 - * Many efficient and provable algorithms.

Set membership (**IN**, **NOT IN**)

- **Query 2:** Find all student names who take CS classes.

Idea: Find the set of sids that take CS classes first. Then check whether any student's id belong to that set or not.

- **IN** is a set membership operator
 - * (**a IN R**) is **TRUE** if **a** appears in **R**

Q: Can we write the same query without subqueries?

Q: Are the above two queries equivalent?

Q: Why we care about duplicates so much?

- **Query 3:** Find the names of students who take no CS classes

Q: Can we rewrite it without subqueries?

Set comparison operator ($> \text{ALL}$, $< \text{SOME}$, \dots)

- **Query 4:** Find the ids of students whose GPA is greater than all students of age 18 or less

– ALL is the universal quantifier \forall

- **Query 5:** Find the IDs of students whose GPA is better than at least one other student of age ≤ 18

– SOME is the existential quantifier \exists

Other Set comparison operators: $> \text{ALL}$, $\leq \text{SOME}$, $= \text{SOME}$, \dots , etc.

– $(< > \text{ALL}) \equiv (\text{NOT IN})$, $(= \text{SOME}) \equiv \text{IN}$

Correlated subqueries

- **Query 6:** Find the names of the students who take any class

- **EXISTS:** WHERE EXISTS(SELECT ... FROM ... WHERE)
 - * True if SELECT .. FROM .. WHERE returns at least one tuple
- **Correlated subquery interpretation:**
 - * Outer query looks at one tuple at a time and binds the tuple to S
 - * For each S, we execute the inner query and check the condition
 - * This is just interpretation. *DBMS executes it more efficiently but get the same result* (but not necessarily MySQL).

Subqueries in FROM clause

- Can be used like a regular relation
- **Example:**

```
SELECT name
FROM (SELECT name, age FROM Student) S
WHERE age > 17
```

 - A subquery inside FROM **MUST** be renamed
 - Student names with age > 17

Common Table Expression

- Introduced in SQL1999
- Similar to subqueries in FROM, but makes it easier to reuse query results
- Syntax:


```
WITH alias AS (query)
SELECT ...
```
- **Example:**

```
WITH S AS (SELECT name, age FROM Student)
SELECT name FROM S WHERE age > 17
```

- **Q:** Do subqueries make SQL more expressive than relational algebra?

Aggregates

- The operators so far check the condition “tuple-by-tuple”
- They never “summarize” multiple tuples into one.
For example, ‘SUM’, ‘AVG’ of GPA is not possible.
- Aggregate function (aggregate diagram)



- **Query 7:** Find the average GPA

- Common aggregate functions: SUM, AVG, COUNT, MIN, MAX on single attribute or COUNT(*).

Problems of Duplicates

- **Query 8:** The number of students taking CS classes
- **Query 9:** The average GPA of the students taking CS classes

GROUP BY clause

- Sometimes, we want to get separate statistics for each group of tuples

Example:

Age	AVG(GPA)
17	3.7
19	2.1
20	3.1

But AVG() takes average over *all* tuples.

- **Query 10:** Find the average GPA for each age group

Q: Is the following query meaningful?

```
SELECT sid, age, AVG(GPA)
FROM Student
GROUP BY age
```

– SELECT can have only attributes that have a single value in each group or *aggregates*

- **Query 11:** Find the number of classes each student is taking

Q: What about the students who take no classes?

Comments: We will learn about outer join that can address this issue later.

HAVING clause

- **Query 12:** Find students who take two or more classes

– Conditions on aggregates should appear in the HAVING clause.

Q: Can we rewrite the query without HAVING clause?

– In general, we can rewrite a query not to have a HAVING clause.