CS143: SQL Query (3)

Book Chapters

(5th) Chapter 3.6-8, 3.11, 4.7-8
(6th) Chapter 3.6, 3.9, 4.1, 5.4-5
(7th) Chapter 3.6, 3.9, 4.1, 5.4-5

Things to Learn

- Window function

- Case function

- ORDER BY and FETCH FIRST

- SQL data modifications

- Null and three-valued logic

- Outer join

- Bag semantics

- SQL expressive power

Window Function

- **Query 1:** Per each result row, return a student's name, their GPA, and the overall GPA average

  - Q: Will this work?

    ```
    SELECT name, GPA, AVG(GPA) FROM Student
    ```

- *Window function*:

  - Syntax: `FTN() OVER()`
    * Append `OVER()` to convert an aggregate function to a window function

– Introduced in SQL 2003

– Aggregate function merges all input tuples into a *single* output tuple

– Window function generates *one output tuple per each input tuple*, but the function is computed over all input tuples

- PARTITION BY:

  – **Query 2:** Per each result row, return a student's name, their GPA, and the average GPA within the student's age group

  – `OVER(PARTITION BY attr)`

  – With `PARTITION BY`, window function is applied only within the same partition

Case Function

- Limited support of if-then-else

  – Return different values depending on conditions

- Syntax: `CASE`

```
            WHEN <condition> THEN <expr>
            WHEN <contidion> THEN <expr>
            ELSE <expr>
    END
```

- Can be used anywhere a column name can be referenced

  – SELECT, WHERE, GROUP BY, ...

- **Query 3:** Average GPA of the child vs adult group

- Q: What if we want to show "child" and "adult" as part of the output?

- Q: What if we want to return two columns, "childGPA" and "adultGPA"?

ORDER BY clause

- Sometimes we may want to display tuples in a certain order. For example order all students by their GPA

- ```
SELECT sid, GPA
FROM Student
ORDER BY GPA DESC, sid ASC
```
    - All students and GPAs, in the descending order of their GPAs and the ascending order of sids. Default is ASC if omitted.
    - Does not change SQL semantics. Just makes the display easier to look at and understand

# FETCH FIRST clause

- **Query 4:** Top-3 students ordered by GPA

    - Sometimes, we just want a few rows from the result. Is there a way to limit result size?

- SQL 2008 Syntax: `[ OFFSET ⟨offset⟩ ROWS ] FETCH FIRST ⟨count⟩ ROWS ONLY`

    - From the result, skip first *offset* rows and return the subsequent *count* rows
    - Unfortunately, this was standardized only in SQL 2008. Many systems use their own syntax, including MySQL.

- Variations:

  - MySQL: `LIMIT` $\langle count \rangle$ `OFFSET` $\langle offset \rangle$
  - Oracle used to use `rownum`, DB2 used to use `SELECT TOP`, but they both support `FETCH FIRST` now
  - MS SQL server requires `ORDER BY` clause *and* `OFFSET` to use `FETCH FIRST`

## General SQL SELECT statement

- ```
  SELECT attributes, aggregates
  FROM relations
  WHERE conditions
  GROUP BY attributes
  HAVING conditions on aggregates
  ORDER BY attributes, aggregates
  FETCH FIRST n ROWS ONLY
  ```

- Evaluation order: FROM → WHERE → GROUP BY → HAVING → ORDER BY → FETCH FIRST → SELECT

Data Modification in SQL (INSERT/DELETE/UPDATE)

- **Insertion**: INSERT INTO *Relation Tuples*

    - Q: Insert tuple (301, CS, 201, 01) to Enroll?

    - Q: Populate Honors table with students of GPA > 3.7?

- **Deletion**: DELETE FROM *R* WHERE *Condition*

    - Q: Delete all students who are not taking classes

- **Update**: Update *R*
            SET *A1 = V1, A2 = V2, ..., An = Vn*
            WHERE *Condition*

    - Q: Increase all CS course numbers by 100

More Advanced SQL

We now go over a bit more esoteric yet important details of SQL

NULL and Three-valued logic

- **Arithmetic operators and comparison**

  Q: `SELECT name`
    `FROM Student`
    `WHERE GPA * 100/4 > 90`
    What should we do if GPA is `NULL`?

  – **Q:** What should be the value for `GPA * 100/4`?

  – Rule: Arithmatic operators with `NULL` input returns `NULL`

  – **Q:** What should be `NULL > 90`?

  – Rule: Arithmatic comparison with `NULL` value return `Unknown`
    * SQL is **Three-valued logic: True, False, Unknown**
    * SQL returns only `True` tuples
    * `GPA * 100/4 > 90` does not return a tuple if GPA is `NULL`

- **Three-valued logic**

  – **Q:** `GPA > 3.7 AND age > 18`. What if GPA is `NULL` and age < 18?

  – **Q:** `GPA > 3.7 OR age > 18`. What if GPA is `NULL` and age < 18?

- Truth table
  * AND: U AND T = U, U AND F = F, U AND U = U
  * OR: U OR T = T, U OR F = U, U OR U = U
- `NOT Unknwon = Unknown`. It's not known
- **SQL returns only `True` tuples**

- **Aggregates**

  - **Q:**

ID	GPA
1	3.0
2	3.6
3	2.4
4	NULL

    `SELECT AVG(GPA)`
    `FROM Student`
    What should be the result?
    What about `COUNT(*)`? `COUNT(GPA)`?

  - Rule: Aggregates are computed ignoring `NULL` value, except `COUNT(*)`.
    * Too much information is lost otherwise.
    * `COUNT(*)` considers a `NULL` tuple as a valid tuple
    * When the input to an aggregate is empty, COUNT returns 0; all others return NULL.

- **Set operators** $(\cup, \cap, -)$

  - **Q:** What should be {2.4, 3.0, `NULL`} $\cup$ {3.6, `NULL`}?

  - Rule: `NULL` is treated like other values in set operators

- **Checking `NULL`**

  - IS NULL or `IS NOT NULL` to check if the value is null.

- `COALESCE()` **function**

  - Return first non-NULL value in the list
  - Example: `COALESCE(phone, email, addr)`

7

OUTER join

- Q: How many classes does each student take?

  - **Q:** What about student 208, Esther? What should we print? What is the problem?

  - **Q:** Anyway to preserve dangling tuples?

- `OUTER JOIN` operator in `FROM` clause:
  - R `LEFT` OUTER JOIN S ON R.A = S.A
    * Keep all dangling tuples from R by padding S attributes with NULL.
  - R `RIGHT` OUTER JOIN S ON R.A = S.A
    * keep all dangling tuples from S by padding R attributes with NULL
  - R `FULL` OUTER JOIN S ON R.A = S.A
    * keep all dangling tuples both from R and S with appropriate padding

- **Q:** How to rewrite the above query to include Esther?

SQL and bag semantics

- What is a bag (multiset)?

  – A set with duplicate elements
  – Order does not matter
  – **Example:** $\{a, a, b, c\} = \{a, c, b, a\} \neq \{a, b, c\}$

- **SQL and bag semantics**

  – Default SQL statements are based on bag semantics
    * We already learned the bag semantics
    * **Except set operators** (UNION, INTERSECT, EXCEPT), which use set semantics
  – We can enforce set semantics by using DISTINCT keyword

- **Bag semantics for set operators**

  – UNION ALL, INTERSECT ALL, EXCEPT ALL
    * MySQL supports only UNION ALL

  – **Q:** $\{a, a, b\} \cup \{a, b, c\}$?

  – **Q:** $\{a, a, a, b, c\} \cap \{a, a, b\}$?

  – **Q:** $\{a, a, b, b\} - \{a, b, b, c\}$?

- **What rules still hold for Bag?**

  – **Q:** Under bag semantics, $R \cup S = S \cup R$? $R \cap S = S \cap R$?
    $R \cap (S \cup T) = (R \cap S) \cup (R \cap T)$?
    * Under bag semantics, some rules still hold, some do not
    * Consider, $R = \{a\}, S = \{a\}, T = \{a\}$ to check the distributive rule.

9

Expressive power of SQL

- **Example:** All ancestors

child	parent
Susan	John
John	James
James	Elaine
...	...

  - **Q:** Can we find all ancestors of Susan using SQL?

- **Example:** All reachable destination

city 1	city 2
A	B
B	D
A	C
E	F
G	H
...	...

  - **Q:** Find all cities reachable from A?

- **Comments:** SQL92 does not support "recursion" and thus cannot compute the *transitive closure.*

  - Recursion is supported in SQL1999.
  - WITH RECURSIVE R(A1, A2) AS ...

    ```
    WITH RECURSIVE Ancestor(child, ancestor) AS (
       (SELECT child, parent AS ancestor FROM Parent)
        UNION
       (SELECT P.child, A.ancestor
        FROM Parent P, Ancestor A
        WHERE P.parent = A.child) )
    SELECT * FROM Ancestor
    ```

  - MySQL introduced support for recursive common table expression in v8.0